

To prediction Crop yield

Importing Important Libraries

```
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error as mse
import warnings
warnings.filterwarnings("ignore")
import matplotlib.pyplot as plt
import numpy as np
```

```
z = pd.read_csv(r"C:\Users\skj_h\OneDrive\Desktop\crop yield
prediction dataset\train.csv")
```

	id	Year	State	Crop_Type	Rainfall	Soil_Type	Irrigation_Area	\
0	1	2019	Punjab	Wheat	578.6	Loamy	3515.2	
1	2	2018	Punjab	Wheat	598.3	Loamy	3499.3	
2	3	2017	Punjab	Wheat	493.0	Loamy	3467.7	
3	4	2016	Punjab	Wheat	426.7	Loamy	3474.6	
4	5	2015	Punjab	Wheat	546.9	Loamy	3474.7	
5	6	2014	Punjab	Wheat	384.9	Loamy	3474.7	
6	7	2013	Punjab	Wheat	619.7	Loamy	3488.1	
7	8	2011	Punjab	Wheat	218.9	Loamy	3466.9	
8	9	2010	Punjab	Wheat	472.1	Loamy	3474.8	
9	10	2009	Punjab	Wheat	384.9	Loamy	3474.8	
10	11	2008	Punjab	Wheat	529.2	Loamy	3437.9	
11	12	2007	Punjab	Wheat	438.0	Loamy	3406.9	
12	13	2006	Punjab	Wheat	418.3	Loamy	3404.8	
13	14	2005	Punjab	Wheat	565.9	Loamy	3410.5	
14	15	2004	Punjab	Wheat	375.2	Loamy	3381.7	
15	16	2003	Punjab	Wheat	459.5	Loamy	3311.6	
16	17	2002	Punjab	Wheat	314.5	Loamy	3353.5	
17	18	2001	Punjab	Wheat	662.8	Loamy	3333.6	
18	19	2000	Punjab	Wheat	391.9	Loamy	3284.3	
19	20	2021	Punjab	Rice	556.9	alluvial	3229.5	
20	21	2020	Punjab	Rice	602.6	alluvial	3118.8	
21	22	2016	Punjab	Rice	426.7	alluvial	2961.4	
22	23	2015	Punjab	Rice	546.9	alluvial	2838.3	
23	24	2014	Punjab	Rice	384.9	alluvial	2838.3	
24	25	2013	Punjab	Rice	619.7	alluvial	2837.6	
25	26	2011	Punjab	Rice	218.9	alluvial	2814.2	
26	27	2010	Punjab	Rice	472.1	alluvial	2721.8	
27	28	2009	Punjab	Rice	384.9	alluvial	2721.8	
28	29	2008	Punjab	Rice	529.2	alluvial	2592.2	
29	30	2007	Punjab	Rice	438.0	alluvial	2602.4	
30	31	2006	Punjab	Rice	418.3	alluvial	2639.9	

31	32	2005	Punjab	Rice	565.9	alluvial	2632.3
32	33	2004	Punjab	Rice	375.2	alluvial	2599.6
33	34	2003	Punjab	Rice	459.5	alluvial	2515.7
34	35	2002	Punjab	Rice	314.5	alluvial	2471.0
35	36	2001	Punjab	Rice	662.8	alluvial	2590.3
36	37	2000	Punjab	Rice	391.9	alluvial	2584.7
37	38	2021	Punjab	Bajra	556.9	Loamy	3.9
38	39	2020	Punjab	Bajra	602.6	Loamy	2.0
39	40	2019	Punjab	Bajra	578.6	Loamy	1.9
40	41	2018	Punjab	Bajra	598.3	Loamy	2.8
41	42	2017	Punjab	Bajra	493.0	Loamy	3.1
42	43	2016	Punjab	Bajra	426.7	Loamy	1.9
43	44	2015	Punjab	Bajra	546.9	Loamy	1.2
44	45	2010	Punjab	Bajra	472.1	Loamy	4.9
45	46	2009	Punjab	Bajra	384.9	Loamy	4.9
46	47	2008	Punjab	Bajra	529.2	Loamy	3.5
47	48	2007	Punjab	Bajra	438.0	Loamy	4.2
48	49	2006	Punjab	Bajra	418.3	Loamy	5.2
49	50	2005	Punjab	Bajra	565.9	Loamy	5.6
50	51	2004	Punjab	Bajra	375.2	Loamy	7.2
51	52	2003	Punjab	Bajra	459.5	Loamy	6.1
52	53	2002	Punjab	Bajra	314.5	Loamy	7.6
53	54	2001	Punjab	Bajra	662.8	Loamy	5.4
54	55	2000	Punjab	Bajra	391.9	Loamy	4.6

Crop_Yield (kg/ha)

0	5188
1	5077
2	5046
3	4583
4	4304
5	5017
6	4724
7	4693
8	4307
9	4462
10	4507
11	4210
12	4179
13	4221
14	4207
15	4200
16	4532
17	4563
18	4696
19	4443
20	4034
21	3974
22	3838

23	3952
24	3998
25	3828
26	4010
27	4022
28	4019
29	3868
30	3858
31	3943
32	3694
33	3510
34	3545
35	3506
36	3347
37	40
38	635
39	583
40	597
41	580
42	0
43	0
44	1495
45	1055
46	950
47	977
48	1045
49	978
50	993
51	810
52	929
53	893
54	703

```
z.isnull().sum()
```

id	0
Year	0
State	0
Crop_Type	0
Rainfall	0
Soil_Type	0
Irrigation_Area	0
Crop_Yield (kg/ha)	0

dtype: int64

```
z.shape
```

```
(55, 8)
```

```
z.size
```

440

```
z.describe()
```

	id	Year	Rainfall	Irrigation_Area	Crop_Yield (kg/ha)
count	55.000000	55.000000	55.000000	55.000000	55.000000
mean	28.00000	2009.527273	473.881818	2082.207273	3079.418182
std	16.02082	6.394021	106.836760	1495.190498	1706.608372
min	1.00000	2000.000000	218.900000	1.200000	0.000000
25%	14.50000	2004.000000	391.900000	5.500000	985.500000
50%	28.00000	2009.000000	459.500000	2721.800000	3943.000000
75%	41.50000	2015.000000	561.400000	3393.250000	4305.500000
max	55.00000	2021.000000	662.800000	3515.200000	5188.000000

```
z.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 55 entries, 0 to 54
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	id	55 non-null	int64
1	Year	55 non-null	int64
2	State	55 non-null	object
3	Crop_Type	55 non-null	object
4	Rainfall	55 non-null	float64
5	Soil_Type	55 non-null	object
6	Irrigation_Area	55 non-null	float64
7	Crop_Yield (kg/ha)	55 non-null	int64


```
dtypes: float64(2), int64(3), object(3)
```


```
memory usage: 3.6+ KB
```


```
z["id"] = z["id"].astype(str)
```

Data Analysis


```
for i in z.columns:
    if(z[i].dtype == "object") and ( i != "id"):
        plt.figure(figsize = (10, 10))
        sns.barplot(x = z[i], y = z["Crop_Yield (kg/ha)"], data = z,
hue = z[i])
```










```
sns.lineplot(x = z["Year"], y = z["Crop_Yield (kg/ha)"], data = z)  
<Axes: xlabel='Year', ylabel='Crop_Yield (kg/ha)'>
```




```
for i in z.columns:
    if(z[i].dtype != "object"):
        plt.figure(figsize = (10, 10))
        sns.scatterplot(x = z[i], y = z["Crop_Yield (kg/ha)"], data =
z)
```

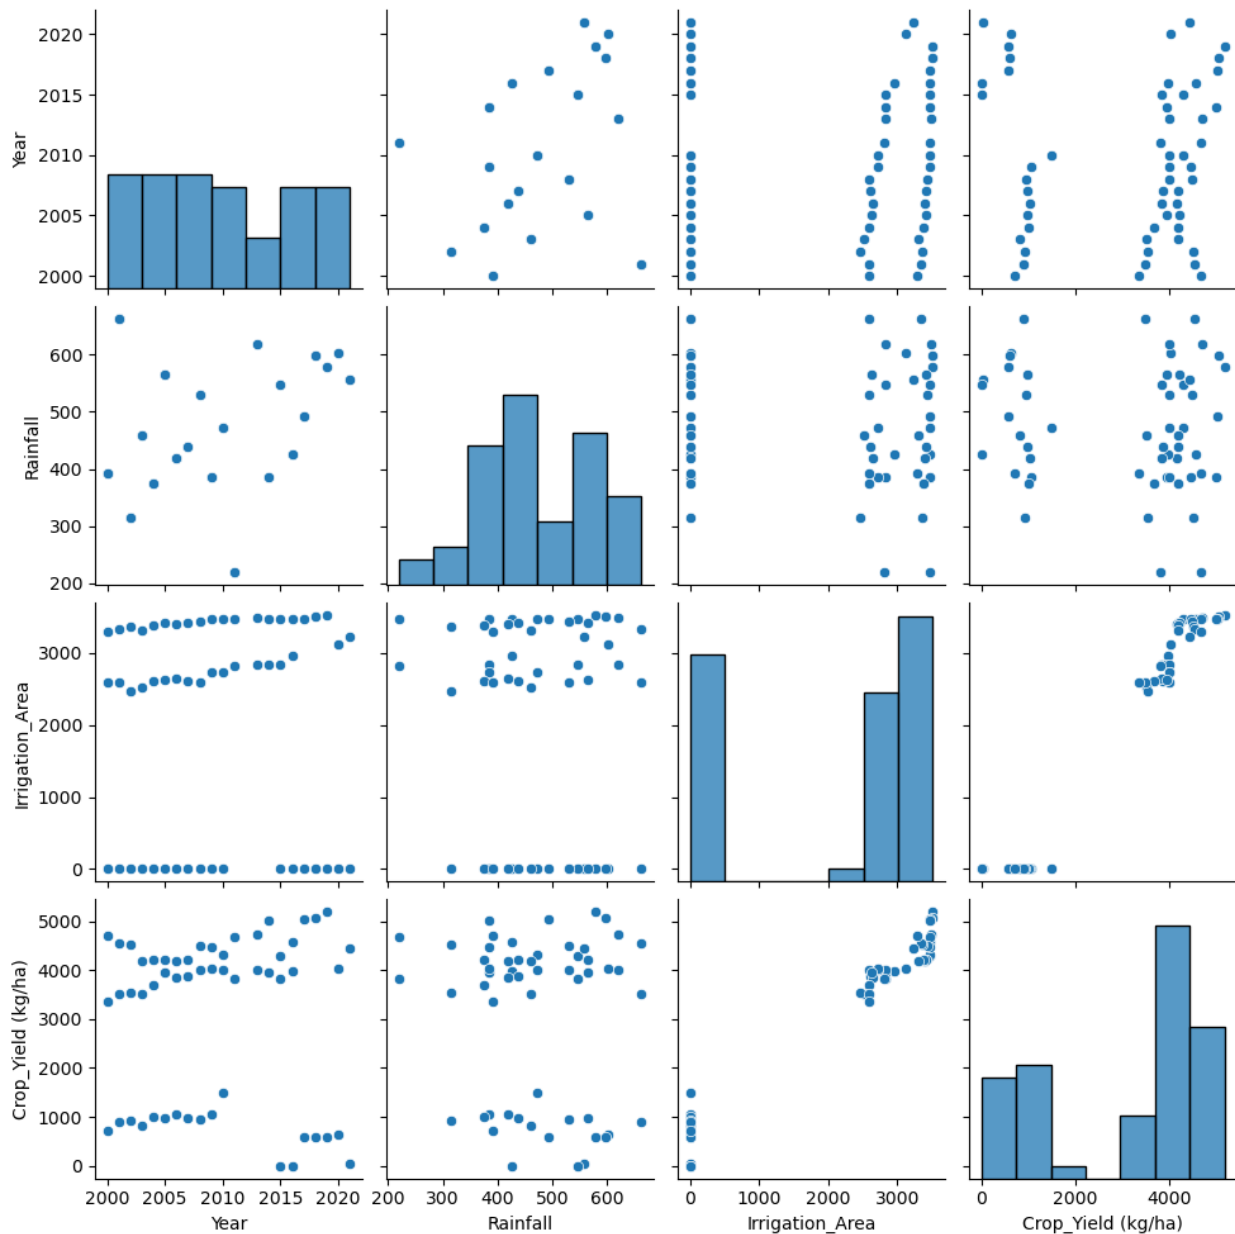






```
sns.pairplot(z)
```

```
<seaborn.axisgrid.PairGrid at 0x1cb95192a80>
```



```
z.columns
```

```
Index(['id', 'Year', 'State', 'Crop_Type', 'Rainfall', 'Soil_Type',  
      'Irrigation_Area', 'Crop_Yield (kg/ha)'],  
      dtype='object')
```

```
pd.DataFrame(z.groupby(["State", "Crop_Type", "Soil_Type"]  
["Crop_Yield (kg/ha)"].sum()))
```

State	Crop_Type	Soil_Type	Crop_Yield (kg/ha)
Punjab	Bajra	Loamy	13263

Rice	alluvial	69389
Wheat	Loamy	86716

```

b = z.copy()
for i in b.columns:
    if(b[i].dtype == "object"):
        b.drop([i], axis = 1, inplace = True)
b

```

	Year	Rainfall	Irrigation_Area	Crop_Yield (kg/ha)
0	2019	578.6	3515.2	5188
1	2018	598.3	3499.3	5077
2	2017	493.0	3467.7	5046
3	2016	426.7	3474.6	4583
4	2015	546.9	3474.7	4304
5	2014	384.9	3474.7	5017
6	2013	619.7	3488.1	4724
7	2011	218.9	3466.9	4693
8	2010	472.1	3474.8	4307
9	2009	384.9	3474.8	4462
10	2008	529.2	3437.9	4507
11	2007	438.0	3406.9	4210
12	2006	418.3	3404.8	4179
13	2005	565.9	3410.5	4221
14	2004	375.2	3381.7	4207
15	2003	459.5	3311.6	4200
16	2002	314.5	3353.5	4532
17	2001	662.8	3333.6	4563
18	2000	391.9	3284.3	4696
19	2021	556.9	3229.5	4443
20	2020	602.6	3118.8	4034
21	2016	426.7	2961.4	3974
22	2015	546.9	2838.3	3838
23	2014	384.9	2838.3	3952
24	2013	619.7	2837.6	3998
25	2011	218.9	2814.2	3828
26	2010	472.1	2721.8	4010
27	2009	384.9	2721.8	4022
28	2008	529.2	2592.2	4019
29	2007	438.0	2602.4	3868
30	2006	418.3	2639.9	3858
31	2005	565.9	2632.3	3943
32	2004	375.2	2599.6	3694
33	2003	459.5	2515.7	3510
34	2002	314.5	2471.0	3545
35	2001	662.8	2590.3	3506
36	2000	391.9	2584.7	3347
37	2021	556.9	3.9	40
38	2020	602.6	2.0	635
39	2019	578.6	1.9	583


40	2018	598.3	2.8	597
41	2017	493.0	3.1	580
42	2016	426.7	1.9	0
43	2015	546.9	1.2	0
44	2010	472.1	4.9	1495
45	2009	384.9	4.9	1055
46	2008	529.2	3.5	950
47	2007	438.0	4.2	977
48	2006	418.3	5.2	1045
49	2005	565.9	5.6	978
50	2004	375.2	7.2	993
51	2003	459.5	6.1	810
52	2002	314.5	7.6	929
53	2001	662.8	5.4	893
54	2000	391.9	4.6	703

```
b.corr()
```

	Year	Rainfall	Irrigation_Area	Crop_Yield
(kg/ha)				
Year	1.000000	0.304973	-0.001326	-
0.026250				
Rainfall	0.304973	1.000000	-0.085025	-
0.092148				
Irrigation_Area	-0.001326	-0.085025	1.000000	
0.984287				
Crop_Yield (kg/ha)	-0.026250	-0.092148	0.984287	
1.000000				

```
sns.heatmap(b.corr(), annot = True)
```

```
<Axes: >
```




```
b.corr()["Crop_Yield (kg/ha)"].sort_values(ascending = False)
```

```
Crop_Yield (kg/ha)    1.000000
Irrigation_Area      0.984287
Year                 -0.026250
Rainfall             -0.092148
Name: Crop_Yield (kg/ha), dtype: float64
```

Regression analysis

```
sns.regplot(x = z["Irrigation_Area"], y = z["Crop_Yield (kg/ha)"],
data = z, line_kws = {"color" : "red"})
<Axes: xlabel='Irrigation_Area', ylabel='Crop_Yield (kg/ha)'>
```

Model selection

```
X = b["Irrigation_Area"]
Y = b["Crop_Yield (kg/ha)"]

x_train, x_test, y_train, y_test = train_test_split(X, Y, train_size =
0.7, test_size = 0.3, random_state = 100)
```

Reshaping x_train

```
x_train = np.array(x_train).reshape(-1, 1)
```

Reshaping y_train

```
y_train = np.array(y_train).reshape(-1, 1)
```

Training model using train dataset

```
n = LinearRegression()
n.fit(x_train, y_train)

LinearRegression()
```

Evaluation of Training dataset

```
y_predict_train = n.predict(x_train)
r2_train = r2_score(y_true = y_train, y_pred = y_predict_train)

round(r2_train, 2)*100

97.0

mse_train = mse(y_true = y_train, y_pred = y_predict_train)
rmse_train = np.sqrt(mse_train)
rmse_train

294.06160727273607
```

Reshaping x_test

```
x_test = np.array(x_test).reshape(-1, 1)
```

Reshaping y_test

```
y_test = np.array(y_test).reshape(-1, 1)
```

Evaluation of Testing dataset

```
y_predict_test = n.predict(x_test)
r2_test = r2_score(y_true = y_test, y_pred = y_predict_test)

round(r2_test, 2)*100

97.0

mse_test = mse(y_true = y_test, y_pred = y_predict_test)
rmse_test = np.sqrt(mse_test)
rmse_test

320.86173720697604
```

Residual analysis for training dataset


```
res_train = y_train - y_predict_train
res_train

array([[ -202.3891355 ],
       [-781.07631834],
       [-106.82293029],
        [ 24.3288216 ],
        [ 136.99624114],
        [ 121.16413507],
```

```
[ 141.68780015],  
[ -29.77273187],  
[ 167.17325878],  
[ 130.88312969],  
[ -44.81163671],  
[-289.13345846],  
[ 469.71807177],  
[-276.36917647],  
[-324.89757696],  
[ 265.31418937],  
[ 710.64163876],  
[ -46.13847648],  
[ 241.93949669],  
[-118.4785779 ],  
[  84.22535659],  
[-259.69645021],  
[ 192.87582875],  
[ 108.09463161],  
[-780.31050833],  
[-271.62584665],  
[ 270.64163876],  
[ 473.28892363],  
[  62.6929756 ],  
[ 190.91396047],  
[ -33.10945026],  
[ 260.31343447],  
[-201.93544427],  
[  67.86152352],  
[-198.07631834],  
[   2.74022496],  
[ 114.62733353],  
[-273.4785779 ]])
```

```
sns.distplot(res_train, kde = True)
```

```
<Axes: ylabel='Density'>
```



Residual analysis for testing dataset


```
res_test = y_test - y_predict_test
```

```
res_test
```

```
array([[ 404.09844254],  
       [ 284.22846915],  
       [  71.00273673],  
       [-296.19500698],  
       [-185.06093121],  
       [ 193.40744877],  
       [ 563.32324442],  
       [ 128.97103192],  
       [ 253.31418937],  
       [ -81.03015695],  
       [ 436.63082353],  
       [-157.0094874 ],  
       [ 206.12540587],  
       [ 323.93114609],  
       [-146.18571977],  
       [ -21.20946357],  
       [-743.26434694]])
```

```
sns.distplot(res_test, kde = True)
```

<Axes: ylabel='Density'>



Recommendations

Crop yield prediction models play a pivotal role in modern agriculture, bridging the gap between data analysis and actionable insights. By leveraging comprehensive datasets and advanced machine learning techniques, stakeholders can drive sustainable growth, improve resource utilization, and secure food supplies for the future. This analysis underscores the transformative potential of data-driven approaches in addressing global agricultural challenges.